

NGON modification proposals

Unstructured grid connectivity

Author: Pierre-Jacques Legay, ONERA

Contact: pierre-jacques.legay@onera.fr

Index table

1 Rationales for the modification proposals.....	1
1.1 Current SIDS description for unstructured grid connectivity.....	1
1.2 Limitations for NGON_N and MIXED.....	3
2 NGON modification proposals.....	4
2.1 Solution description.....	4
2.2 Solution analysis.....	5
3 Example of extension.....	8
4 Implementation note.....	9
5 Conclusions.....	9
6 Annexe : Document modification list.....	9

This CPEX focuses on the NGON representation. The rationale for requiring an extension to CGNS/SIDS for unstructured grid connectivity is detailed in the first part of this document. The second part details the proposal which includes a solution to the problem and an impact analysis of that solution on the SIDS and on the performances.

1 Rationales for the modification proposals

The first section is a reminder of the current SIDS description for unstructured grid connectivity. The following section details the problems induced by this description.

1.1 Current SIDS description for unstructured grid connectivity

The unstructured grid connectivity is stored in the Elements_t of the CGNS/SIDS. As described hereafter, its storage depends on the elements type:

For all element types except MIXED, NGON_n, and NFACE_n, ElementConnectivity contains the list of nodes for each element. If the elements are sorted, then it must first list the connectivity of the boundary elements, then that of the interior elements.

```
ElementConnectivity = Node11, Node21, ... NodeN1,  
                     Node12, Node22, ... NodeN2,  
                     ...  
                     Node1M, Node2M, ... NodeNM
```

where M is the total number of elements (i.e., ElementSize), and N is the number of nodes per element.

ElementDataSize indicates the total size (number of integers) of the array ElementConnectivity. For all element types except MIXED, NGON_n, and NFACE_n, the ElementDataSize is given by:

```
ElementDataSize = ElementSize * NPE[ElementType]
```

where NPE[ElementType] is a function returning the number of nodes for the given ElementType. For example, NPE[HEXA_8]=8.

33 When the section **ElementType** is **MIXED**, the data array **ElementConnectivity**
 34 contains one extra integer per element, to hold each individual element type:
 35 **ElementConnectivity** = **Etype**₁, **Node**₁₁, **Node**₂₁, ... **Node**_{N₁},
 36 **Etype**₂, **Node**₁₂, **Node**₂₂, ... **Node**_{N₂},
 37 ...
 38 **Etype**_M, **Node**_{1M}, **Node**_{2M}, ... **Node**_{N_M}
 39 where again **M** is the total number of elements, and **N_i** is the number of nodes in
 40 element **i**. In the case of **MIXED** element section, **ElementDataSize** is given by:
 41 **ElementDataSize** =sum(**n**=[start, end], **NPE**[**ElementType**_n] + 1)

42 **Arbitrary polyhedral elements may be defined using the NGON_n and**
 43 **NFACE_n element types.** The **NGON_n** element type is used to specify all the faces
 44 in the grid, and the **NFACE_n** element type is then used to define the polyhedral
 45 elements as a collection of these faces.

46 I.e., for **NGON_n**, the data array **ElementConnectivity** contains a list of nodes making
 47 up each face in the grid, with the first value for each face defining the number of
 48 nodes making up that face:

49 **ElementConnectivity** = **Nnodes**₁, **Node**₁₁, **Node**₂₁, ... **Node**_{N₁},
 50 **Nnodes**₂, **Node**₁₂, **Node**₂₂, ... **Node**_{N₂},
 51 ...
 52 **Nnodes**_M, **Node**_{1M}, **Node**_{2M}, ... **Node**_{N_M}

53 where here **M** is the total number of faces, and **N_i** is the number of nodes in face **i**. The
 54 **ElementDataSize** is the total number of nodes defining all the faces, plus one value
 55 per face specifying the number of nodes making up that face.

56 Then for **NFACE_n**, **ElementConnectivity** contains the list of face elements making up
 57 each polyhedral element, with the first value for each polyhedra defining the number
 58 of faces making up that polyhedral element.

59 **ElementConnectivity** = **Nfaces**₁, **Face**₁₁, **Face**₂₁, ... **Face**_{N₁},
 60 **Nfaces**₂, **Face**₁₂, **Face**₂₂, ... **Face**_{N₂},
 61 ...
 62 **Nfaces**_M, **Face**_{1M}, **Face**_{2M}, ... **Face**_{N_M}

63 where now **M** is the total number of polyhedral elements, and **N_i** is the number of faces
 64 in element **i**. The sign of the face number determines its orientation. If the face
 65 number is positive, the face normal is directed outward; if it's negative, the face
 66 normal is directed inward.

67 **ElementDataSize** =sum(**n**=[start, end], **FPP**[**ElementType**_n] + 1)

68 where **FPP**[**ElementType**_n] is a function returning the number of faces per polyhedra.

69 The following figure is set up to ease comprehension and comparison in the following sections:

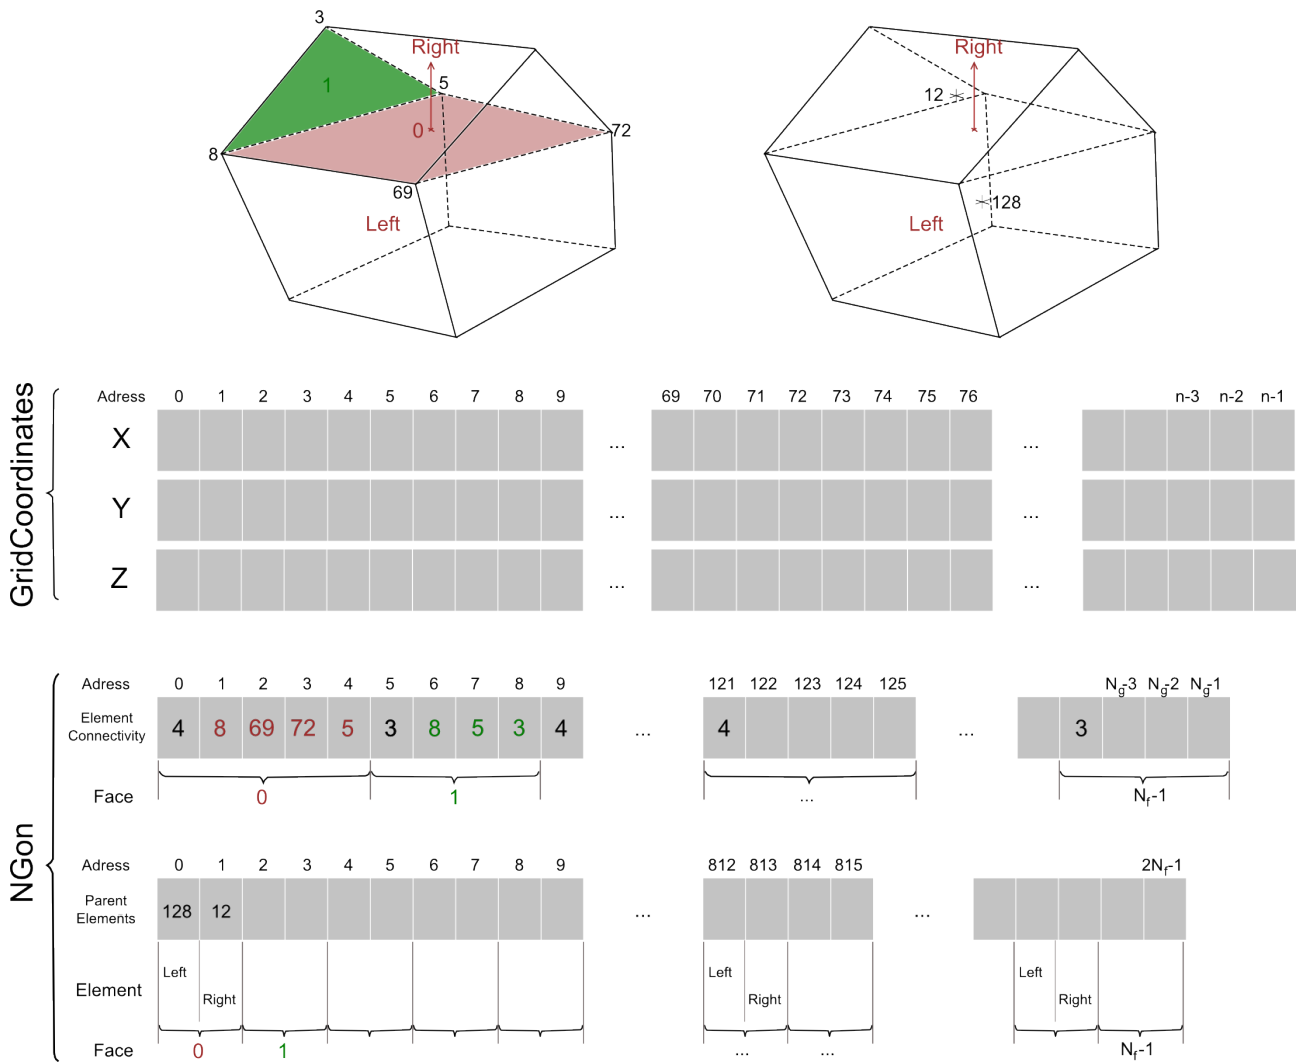


Illustration 1: CGNS/NGON current face based representation.

70 This figure shows an unstructured element using the NGON representation. It describes the physical
71 construction of the CGNS related arrays GridCoordinates, ElementConnectivity, and
72 ParentElements.

73 Notation:

- 74 • N_g : represents a vertex from the current face.
- 75 • N_f : represents a face of the current element.

76 1.2 Limitations for NGON_N and MIXED

77 The ElementConnectivity array can be very large for industrial CFD case. As a direct consequence
78 we should be able to load this array fully or partially on multiple threads.

79 In the above description of the CGNS/NGON face based representation, the ElementConnectivity
80 array mixes two data types which are interdependent :

- 81 • the number of nodes for each face : N_{nodes1}
- 82 • a list of nodes for a face : $Node_{11}, Node_{21}, \dots, Node_{N1}$

83 This implies that we cannot efficiently split this array to be read in parallel.



Illustration 2: ElementConnectivity interlaced data.

84 In parallel, these interlaced data impose the load of the complete ElementConnectivity array on
 85 every processors. This has a significant impact on IO performances.

86 **NB: As described in the SIDS, the representation of MIXED and NFACE elements is identical**
 87 **to the NGON representation. As such, the parallel read suffers from the same performance**
 88 **issues. The issues of the MIXED and NFACE elements shall be addressed in dedicated CPEX.**

89 2 NGON modification proposals

90 This proposal offers to alter the NGON SIDS representation in order to allow an efficient parallel
 91 IO access and to optimize the data representation.

92 2.1 Solution description

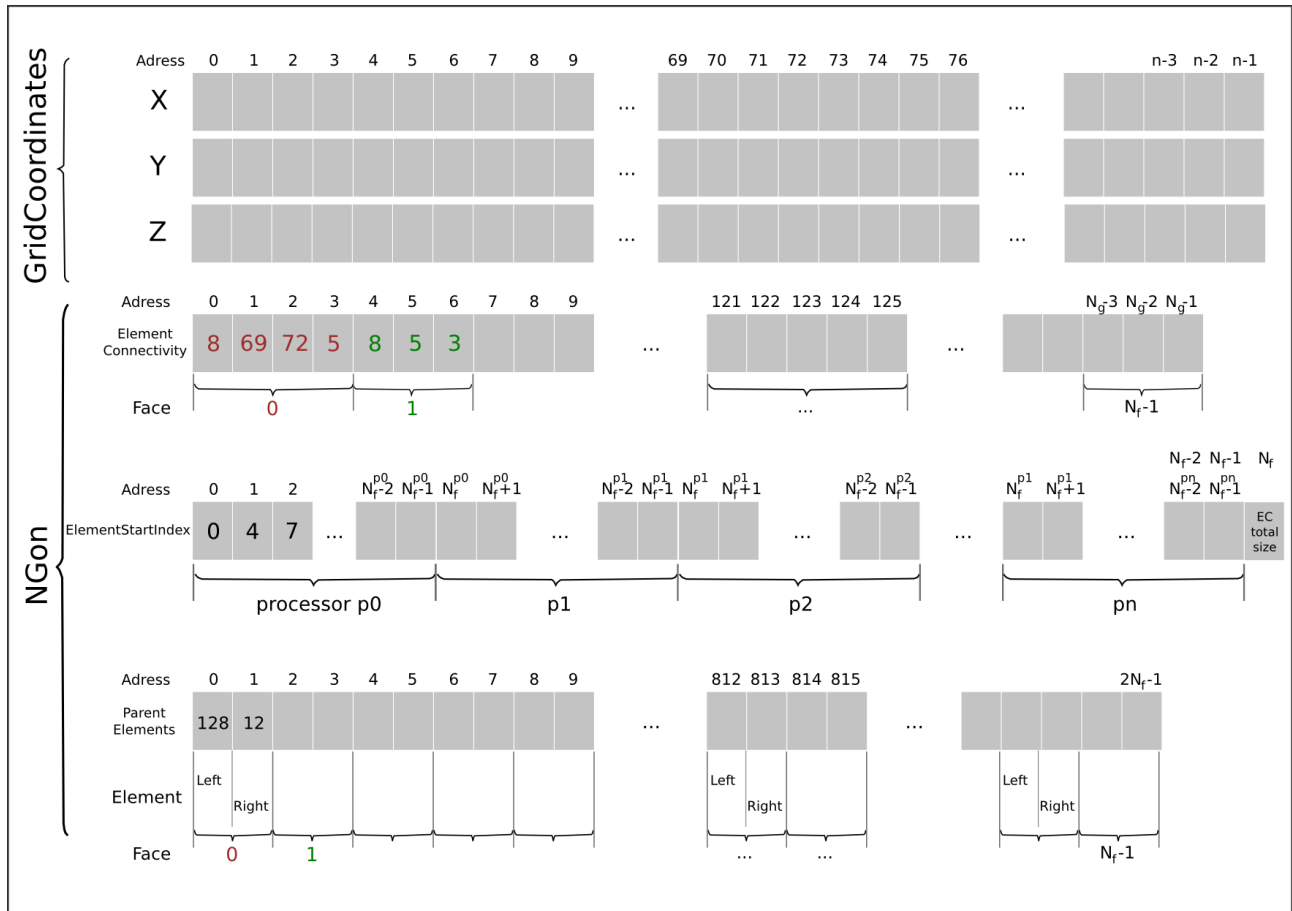


Illustration 3: NGON with new ElementConnectivity array and ElementStartIndex position array.

93 In this representation the ElementConnectivity array is deinterlaced. The following pictures
 94 compares the current standard versus the new deinterlaced ElementConnectivity array:

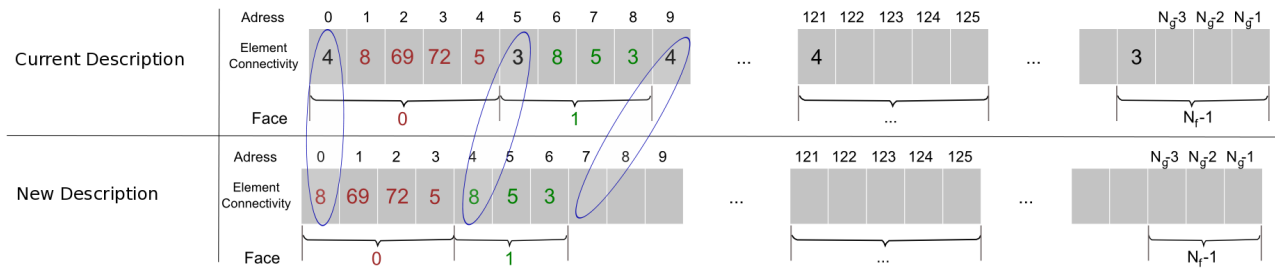


Illustration 4: Current ElementConnectivity vs new ElementConnectivity.

95 The face vertex count has been removed from the connectivity. As a direct consequence the new
 96 array includes a unique data type.

97 Adresses needed to read or analyze the ElementConnectivity array must be obtained from the new
 98 ElementStartIndex array.

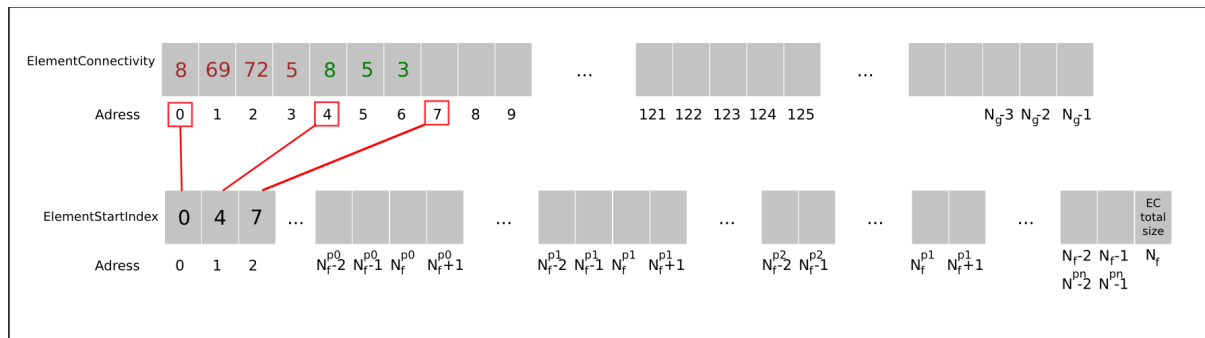


Illustration 5: ElementStartIndex lists the face position in the ElementConnectivity and it's last value indicates the ElementConnectivity total size.

99 This array lists the position in the ElementConnectivity of the first vertex for each face and it's last
 100 value is the ElementConnectivity array size. This CGNS node is of type DataArray_t. Its read can
 101 be distributed between P processors. Its size is N_f+1 with N_f being the number of face of the
 102 ElementConnectivity array.

103 The ElementStartIndex array makes it easy for a process or thread to read a portion of the
 104 ElementConnectivity array.

105 NB: The last value of the ElementStartIndex allows easy access to the last vertex of the last face of
 106 the ElementConnectivity array.

107 2.2 Solution analysis

108 1. This solution modifies the CGNS standard.
 109 It strongly modifies the representation of the ElementConnectivity node and it inserts the
 110 new array ElementStartIndex.
 111 These modifications could be challenging for the current CFD database.

112 2. Data consistency
 113 The data type of both arrays ElementConnectivity and ElementStartIndex are consistent.

3. Optimal data size

This solution does not duplicate data thus the global data size is unchanged and stays optimal.

4. ElementStartIndex incremental index

In this proposal we choose to use the face position in the element connectivity instead of the face number of nodes (ElementStartIndex=[0,4,7,11, ...] instead of [4, 3, 4, ...]).

The rationales are all IO read/write oriented and detailed here after:

- Processor inter-dependency

To access the ElementConnectivity data we need to know the position of the face in the ElementConnectivity array.

→ “Face vertex number” solution

In this configuration we would need to sum over the index array to obtain the face position in the ElementConnectivity array. This means either create a new table or perform the computation as many times as needed.

Another problem is that we need the result of the sum of face vertex of processor P_0 to obtain the location of the first vertex of the first face assigned to processor P_1 .

This behavior generalizes with the need to sum the face vertex of processors P_0, P_1, \dots, P_{M-1} to obtain the location of the first vertex of the first face assigned to processor P_M .

This is not a complex operation but we do not like the idea of imposing a dependency on all previous processors computations.

→ “Face position” solution

In this configuration the ElementStartIndex array can be split on P processors and each processor related part directly gives the location of the face vertex from the ElementConnectivity array.

- Full vs partial load

As stated above we need to know the position of the face in the ElementConnectivity array.

→ “Face vertex number” solution

If we want to access data in processor P_M , then we need to access data on all processors P_0, P_1, \dots, P_{M-1} . This means that we need to load the index array on every processor of inferior rank.

→ “Face position” solution

To access the ElementConnectivity we need the boundary of locations for processor P_M .

These boundaries can be partially loaded from the ElementStartIndex array by getting the first element of processor P_M part and the first element of processor P_{M+1} part. These two integers indicate the section of ElementConnectivity to be loaded on proc P_M (1st element of P_M till 1st element of $P_{M+1}-1$).

NB: The last value of the ElementStartIndex allows to apply the same operator to the last section of the ElementConnectivity array. For the last section, the load is performed from P_M to $P_{M+1}-1$ with P_{M+1} being the ElementConnectivity size.

With this proposal, we need to load only two integers from the ElementStartIndex array on a processor to fully access the ElementConnectivity array. Partial load for CGNS/HDF5 is available in CHLone for example.

- direct access to a face connectivity

For some specific unstructured software, it is interesting to have a simple access to a specific face in the ElementConnectivity array.

This is only allowed with the “Face Position” solution.

- Access to position and face number

Using the incremental index method, it is easy to calculate the “face number of nodes” for

162 element 'i' as `ElementStartIndex[i+1]-ElementStartIndex[i]` is an $O(1)$ calculation. So the
163 selected method gives both data (face position and face number of nodes) in $O(1)$ where the
164 alternate method would give face number of nodes in $O(1)$, but face position in $O(N)$.

165 5. `ElementStartIndex` content

166 Using an incremental index -the face position in the `ElementConnectivity` array- could lead
167 to large numbers in the `ElementStartIndex` array. However, these numbers are limited to the
168 size of the `ElementConnectivity` array as they give access to addresses of that array. We will
169 probably reach the `ElementConnectivity` array limits first.

170 6. `ElementStartIndex` last value

171 The `ElementStartIndex` lists the first vertex positions of each faces in the
172 `ElementConnectivity`. Then we added the `ElementConnectivity` size as last value. The
173 rational behind this is to improve data access when looping over the faces.

174 → Positions without `ElementConnectivity` size

175 In this configuration an iteration loop over NGONs could be written as follow:

```
176 for (i=0; i < NGON; i++) {  
177     end = (I == NGON-1) ? ElementConnectivitySize :ElementStartIndex[i+1];  
178     for (j=ElementStartIndex[i]; j < end; j++) {  
179         entry = ElementConnectivity[j];  
180         // do something here  
181     }  
182 }
```

183 For each NGON we have to check if we are reaching the last NGON to handle the access to
184 the last vertex of the last face.

185 → Positions with `ElementConnectivity` size (current proposal)

186 In this configuration an iteration loop over NGONs could be written as follow:

```
187 for (i=0; i < NGON; i++) {  
188     for (j=ElementStartIndex[i]; j < ElementStartIndex[i+1]; j++) {  
189         entry = ElementConnectivity[j];  
190         // do something here  
191     }  
192 }
```

193 The last value of the `ElementStartIndex` gives a boundary to access the last vertex of the last
194 NGON without condition.




Illustration 6: Simple configuration demonstrating the NGON proposal.

Node Tree

/

CGNSLibraryVersion

Base

Zone

ZoneType

GridCoordinates

NGonElements

ElementRange

ElementConnectivity

ParentElements

ZoneBC

BCInflow.0

BCWall.3

BCWall.4

BCWall.5

BCWall.6

BCOutflow.0

Node Description

Parent Node

/Base/Zone/NGonElements

Node Name

ElementConnectivity

Node Label

DataArray_t

Link Description

Link File

Link Node

Data Description

Data Type

i4

Dimensions

740

Bytes

2960

create

modify

Node Data

4 1 2 5 4
4 10 11 14 13
4 10 11 2 1
4 4 5 14 13
4 1 4 13 10
4 2 5 14 11
4 2 3 6 5
4 11 12 15 14
4 11 12 3 2
4 5 6 15 14
4 1 15 6 3

Line 22 (211)

Node Description

Parent Node

/Base/Zone/NGonElements

Node Name

ElementConnectivity

Node Label

DataArray_t

Link Description

Link File

Link Node

Data Description

Data Type

i4

Dimensions

555

Bytes

2220

create

modify

Node Data

1 2 5 4
10 11 14 13
10 11 2 1
4 5 14 13
1 4 13 10
2 5 14 11
2 3 6 5
11 12 15 14
11 12 3 2
5 6 15 14
1 2 15 6 3

Line 1 (1)

Illustration 7: Previous vs new ElementConnectivity array

Node Tree

/

CGNSLibraryVersion

Base

Zone

ZoneType

GridCoordinates

NGonElements

ElementRange

ElementConnectivity

ParentElements

ElementStartIndex

ZoneBC

Node Description

Parent Node

/Base/Zone/NGonElements

Node Name

ElementStartIndex

Node Label

DataArray_t

Link Description

Link File

Browse

Link Node

Browse

Data Description

Data Type

i4

Dimensions

186

Bytes

740

create

modify

read

clear

delete

Node Data

0 4 8 12 16 20 24 28 32 36
40 44 48 52 56 60 64 68 72 76
80 84 88 92 96 100 104 108 112 116
120 124 128 132 136 140 144 148 152 156
160 164 168 172 176 180 184 188 192 196
200 204 208 212 216 220 224 228 232 236
240 244 248 252 256 260 264 268 272 276
280 284 288 292 296 300 304 308 312 316
320 324 328 332 336 340 344 348 352 356
360 364 368 372 376 380 384 388 392 396
400 404 408 412 416 420 424 428 432 436

Line 11 (101) Values/Line 10

Illustration 8: ElementStartIndex array

196 4 Implementation note

197 This section aims to reflect discussions held during previous CGNS meetings related to the
198 implementation impact of this modification proposal. As detailed above, this proposal strongly
199 modifies the representation of the ElementConnectivity node. As a consequence we should address
200 the backward compatibility issue.

201 The following solutions were discussed in order to insure backward compatibility:

- 202 • CGNSLibraryVersion
203 In softwares, the CGNSLibraryVersion could be used to switch between actual and future
204 NGON representation. This number attached to the standard version will allow any parsing
205 tools to use the correct NGON representation.
- 206 • Conversion utility
207 A dedicated tool can be built to insure the conversion of the current CGNS files to the new
208 NGON representation.

209 These two solutions should allow a smooth transition for all users of the CGNS standard.

210 5 Conclusions

211 This modification proposal for the NGON representation addresses the HPC issue due to the
212 interlaced data representation of the unstructured connectivity description. The above solution
213 optimizes the data representation for parallel IO and has a low impact on the SIDS to deinterlace
214 data.

215 As shown in the first part, MIXED and NFACES elements are equally concerned by the parallel IO
216 access. As such their SIDS representation should be updated accordingly to the solution chosen for
217 NGON. This issue should be addressed in a later CPEX.

218 6 Annexe : Document modification list

219 1. Following Marc Poinot remarks:

- 220 • Rename array from FaceConnectivityPosition to ElementStartIndex (text and figures)
- 221 • Reformulate a sentence describing the ElementStartIndex at line 101
- 222 • Remove remark concerning array size limitation in section 2.1.2.5

223 2. Following Robert Bush suggestion:

- 224 • Added section “Implementation note” to reflect CGNS committee discussions.

225 3. Remove multiple typo.

226 4. Following Gregory Sjaardema feedback:

- 227 • Modify the ElementStartIndex definition to improve the ability to loop on the data array.
228 The new size is N+1 and last index represents the ElementConnectivity total size.
229 This modification is propagated through sections 2 and 3.