

NGON modification proposal

Unstructured grid connectivity

Author: Pierre-Jacques Legay, ONERA

Contact: pierre-jacques.legay@onera.fr

Index table

1 Rationale for the modification proposal.....	1
1.1 Current SIDS description for unstructured grid connectivity.....	1
1.2 Limitations for NGON_N and MIXED.....	3
2 NGON modification proposal.....	4
2.1 Solution description.....	5
2.2 Solution analysis.....	6
3 Example of extension.....	9
4 Implementation note.....	10
5 Conclusions.....	10
6 Appendix: Document modification list.....	10

This CPEX focuses on the NGON representation. The rationale for requiring an extension to CGNS/SIDS for unstructured grid connectivity is detailed in the first part of this document. The second part details the proposal which includes a solution to the problem and an impact analysis of that solution on the SIDS and on the performance.

1 Rationale for the modification proposal

The first section is a reminder of the current SIDS description for unstructured grid connectivity. The following section details the problems induced by this description.

1.1 Current SIDS description for unstructured grid connectivity

The unstructured grid connectivity is stored in the Elements_t of the CGNS/SIDS. As described hereafter, its storage depends on the elements type:

For all element types except MIXED, NGON_n, and NFACE_n, ElementConnectivity contains the list of nodes for each element. If the elements are sorted, then it must first list the connectivity of the boundary elements, then that of the interior elements.

```
ElementConnectivity = Node11, Node21, ... NodeN1,  
                      Node12, Node22, ... NodeN2,  
                      ...  
                      Node1M, Node2M, ... NodeNM
```

where M is the total number of elements (i.e., ElementSize), and N is the number of nodes per element.

ElementDataSize indicates the total size (number of integers) of the array ElementConnectivity. For all element types except MIXED, NGON_n, and NFACE_n, the ElementDataSize is given by:

```
ElementDataSize = ElementSize * NPE[ElementType]
```

where NPE[ElementType] is a function returning the number of nodes for the given ElementType. For example, NPE[HEXA_8]=8.

32 When the section `ElementType` is MIXED, the data array `ElementConnectivity` contains
 33 one extra integer per element, to hold each individual element type:
 34 `ElementConnectivity = Etype1, Node11, Node21, ... NodeN1,`
 35 `Etype2, Node12, Node22, ... NodeN2,`
 36 `...`
 37 `EtypeM, Node1M, Node2M, ... NodeNM`
 38 where again M is the total number of elements, and N_i is the number of nodes in
 39 element i . In the case of MIXED element section, `ElementDataSize` is given by:
 40 `ElementDataSize =sum(n=[start, end], NPE[ElementTypen] + 1)`

41 Arbitrary polyhedral elements may be defined using the `NGON_n` and `NFACE_n`
 42 element types. The `NGON_n` element type is used to specify all the faces in the grid,
 43 and the `NFACE_n` element type is then used to define the polyhedral elements as a
 44 collection of these faces.

45 I.e., for `NGON_n`, the data array `ElementConnectivity` contains a list of nodes making
 46 up each face in the grid, with the first value for each face defining the number of
 47 nodes making up that face:

48 `ElementConnectivity = Nnodes1, Node11, Node21, ... NodeN1,`
 49 `Nnodes2, Node12, Node22, ... NodeN2,`
 50 `...`
 51 `NnodesM, Node1M, Node2M, ... NodeNM`

52 where here M is the total number of faces, and N_i is the number of nodes in face i . The
 53 `ElementDataSize` is the total number of nodes defining all the faces, plus one value
 54 per face specifying the number of nodes making up that face.

55 Then for `NFACE_n`, `ElementConnectivity` contains the list of face elements making up
 56 each polyhedral element, with the first value for each polyhedra defining the number
 57 of faces making up that polyhedral element.

58 `ElementConnectivity = Nfaces1, Face11, Face21, ... FaceN1,`
 59 `Nfaces2, Face12, Face22, ... FaceN2,`
 60 `...`
 61 `NfacesM, Face1M, Face2M, ... FaceNM`

62 where now M is the total number of polyhedral elements, and N_i is the number of faces
 63 in element i . The sign of the face number determines its orientation. If the face
 64 number is positive, the face normal is directed outward; if it's negative, the face
 65 normal is directed inward.

66 `ElementDataSize =sum(n=[start, end], FPP[ElementTypen] + 1)`

67 where `FPP[ElementTypen]` is a function returning the number of faces per polyhedra.

68 The following figure is set up to ease comprehension and comparison in the following sections:

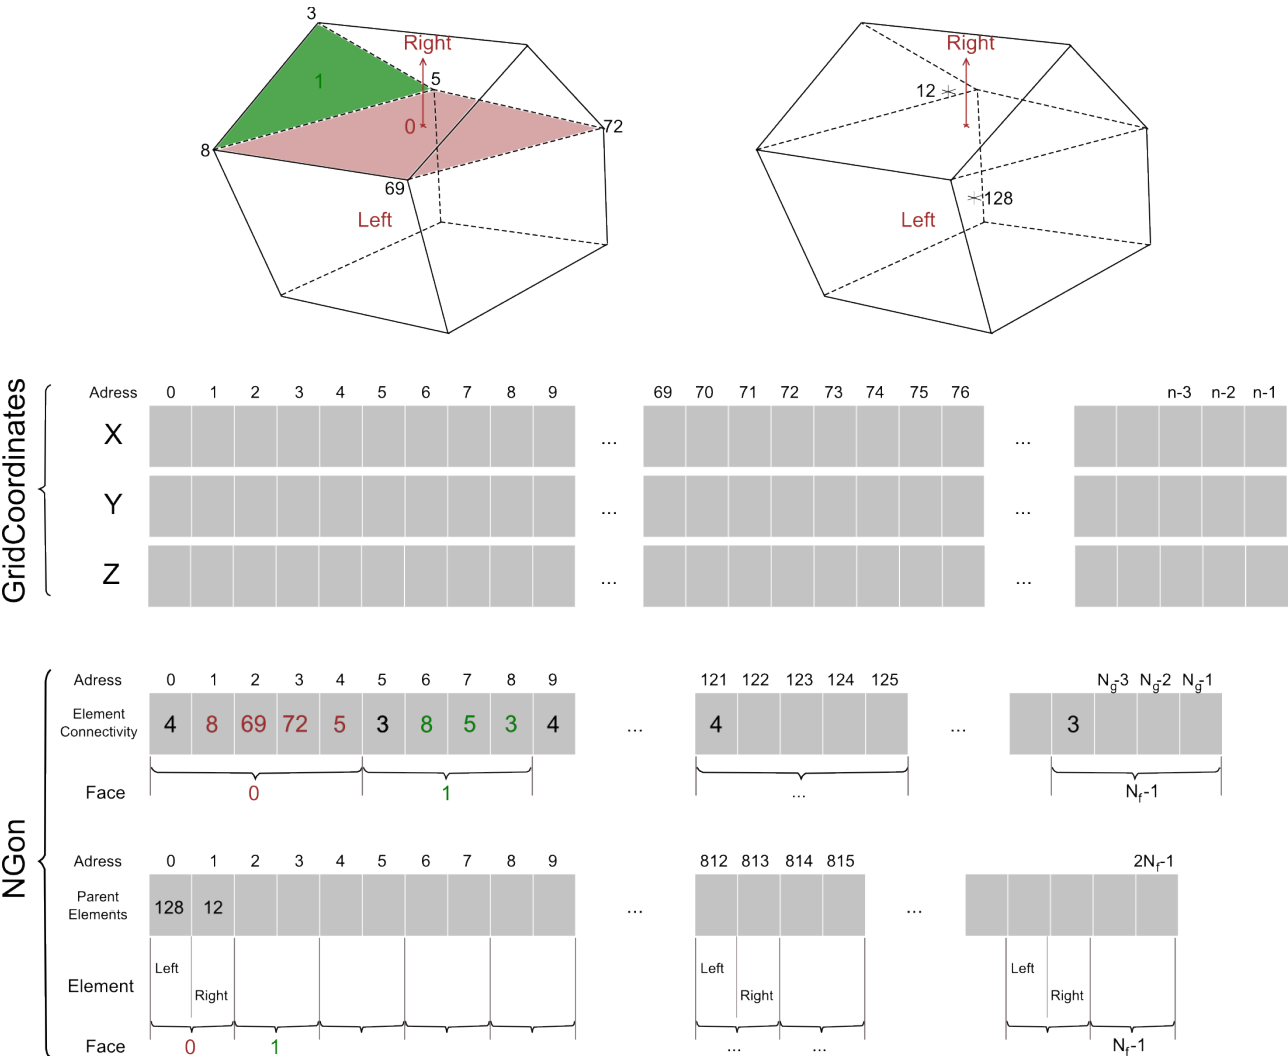


Illustration 1: CGNS/NGON current face based representation.

69 This figure shows an unstructured element using the NGON representation. It describes the physical
70 construction of the CGNS related arrays GridCoordinates, ElementConnectivity, and
71 ParentElements.

72 Notation:

- 73 • N_g : represents a vertex from the current face.
- 74 • N_f : represents a face of the current element.


75 1.2 Limitations for NGON_N and MIXED

76 The ElementConnectivity array can be very large for industrial CFD case. As a direct consequence
77 we should be able to load this array fully or partially on multiple threads.

78 In the above description of the CGNS/NGON face based representation, the ElementConnectivity
79 array mixes two data types which are interdependent :

- 80 • the number of nodes for each face : N_{nodes1}
- 81 • a list of nodes for a face : $Node11, Node21, \dots, NodeN1$

82 This implies that we cannot efficiently split this array to be read in parallel.



85 NB: As described in the SIDS, the representation of MIXED and NFACE elements is identical to
86 the NGON representation. As such, the parallel read suffers from the same performance issues. The
87 issues of the MIXED and NFACE elements shall be addressed in a later CPEX.

89 This proposal modifies the NGON SIDS representation to provide efficient parallel IO access and to
90 optimize the data representation.

The diagram illustrates the mapping of a 3D grid to a 1D array for a GPU-based simulation. It is divided into three main sections: GridCoordinates, NGon, and Element Connectivity.

GridCoordinates: Shows a 3D grid with axes X, Y, and Z. The grid is divided into blocks of size 8x8x8. The addresses are labeled from 0 to 9, 69 to 76, and n-3 to n-1. The grid is mapped to a 1D array.

NGon: Shows the mapping of the grid to a 1D array. The addresses are labeled from 0 to 9, 121 to 125, and N_g-3 to N_g-1 . The grid is mapped to a 1D array.

Element Connectivity: Shows the mapping of the grid to a 1D array. The addresses are labeled from 0 to 9, 121 to 125, and N_g-3 to N_g-1 . The grid is mapped to a 1D array.

Illustration 3: NGON with new ElementConnectivity array and ElementStartIndex position array.

92 In this representation the ElementConnectivity array is deinterlaced. Illustration 4 compares the
 93 current standard versus the new deinterlaced ElementConnectivity array:




Illustration 4: Current ElementConnectivity vs. new ElementConnectivity.

94 The face vertex count has been removed from the connectivity. As a direct consequence the new
 95 array contains a unique data type.

96 Indices needed to read or analyze the ElementConnectivity array are stored in the new
 97 ElementStartIndex array.

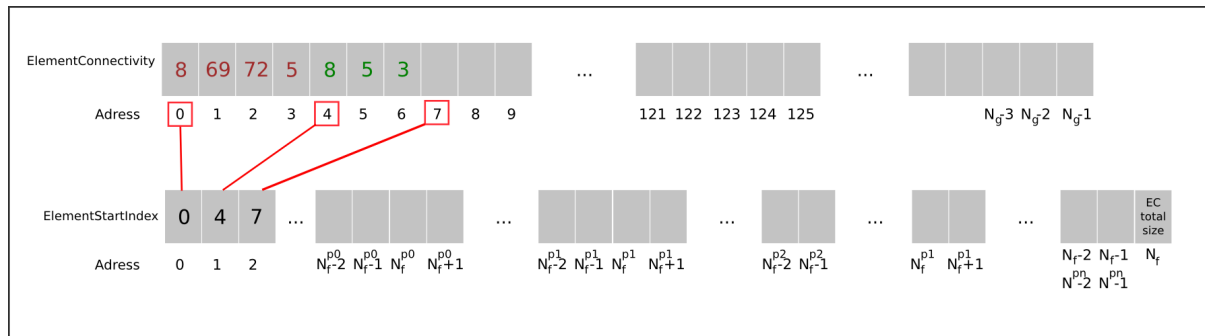


Illustration 5: ElementStartIndex lists the face position in the ElementConnectivity and it's last value indicates the ElementConnectivity total size.

98 This array lists the position in the ElementConnectivity of the first vertex for each face and it's last
 99 value is the ElementConnectivity array size. This CGNS node is of type DataArray_t. Its read can
 100 be distributed between P processors. Its size is $N_f + 1$ with N_f being the number of face in the
 101 ElementConnectivity array. The ElementStartIndex array makes it easy for a process or thread to
 102 read a portion of the ElementConnectivity array.

103 NB: The last value of the ElementStartIndex allows easy access to the last vertex of the last face of
 104 the ElementConnectivity array.

105 2.2 Solution analysis

106 2.2.1 CGNS standard modification

107 It strongly modifies the representation of the ElementConnectivity node and it inserts the
 108 new array ElementStartIndex. These modifications should be reflected in the current CFD
 109 database and CGNS related code to insure the continuity of the computational capability.
 110 See section 4 'Implementation note'.

111 2.2.2 Data consistency

112 The data type of both arrays ElementConnectivity and ElementStartIndex are consistent.

113 2.2.3 Optimal data size

114 This solution does not duplicate data thus the global data size is unchanged and stays
115 optimal.

116 2.2.4 ElementStartIndex -- an incremental index

117 In this proposal we choose to use the face position in the element connectivity instead of the
118 face number of nodes (ElementStartIndex=[0,4,7,11, ...] instead of [4, 3, 4, ...]).
119 The rationales are all IO read/write oriented and detailed here after:

120 1. Processor inter-dependency

121 To access the ElementConnectivity data we need to know the position of the face in the
122 ElementConnectivity array.

123 → “Face vertex number” solution

124 In this configuration we would need to sum over the index array to obtain the face position
125 in the ElementConnectivity array. This means either create a new table or perform the
126 computation as many times as needed.

127 Another problem is that we need the result of the sum of face vertex of processor P_0 to
128 obtain the location of the first vertex of the first face assigned to processor P_1 .

129 This behavior generalizes with the need to sum the face vertex of processors P_0, P_1, \dots, P_{M-1} to
130 obtain the location of the first vertex of the first face assigned to processor P_M .

131 **This is not a complex operation but we do not like the idea of imposing a dependency**
132 **on all previous processors computations.**

133 → “Face position” solution

134 In this configuration the ElementStartIndex array can be split on P processors and each
135 processor related part directly gives the location of the face vertex from the
136 ElementConnectivity array.

137 2. Full vs. partial load

138 As stated above we need to know the position of the face in the ElementConnectivity array.

139 → “Face vertex number” solution

140 If we want to access data in processor P_M , then we need to access data on all processors
141 P_0, P_1, \dots, P_{M-1} . This means that we need to load the index array on every processor of inferior
142 rank.

143 → “Face position” solution

144 To access the ElementConnectivity we need the boundary of locations for processor P_M .

145 These boundaries can be partially loaded from the ElementStartIndex array by getting the
146 first element of processor P_M part and the first element of processor P_{M+1} part. These two
147 integers indicate the section of ElementConnectivity to be loaded on proc P_M (1st element of
148 P_M till 1st element of $P_{M+1}-1$).

149 NB: The last value of the ElementStartIndex allows to apply the same operator to the last
150 section of the ElementConnectivity array. For the last section, the load is performed from P_M
151 to $P_{M+1}-1$ with P_{M+1} being the ElementConnectivity size.

152 **With this proposal, we need to load only two integers from the ElementStartIndex**
153 **array on a processor to fully access the ElementConnectivity array.** Partial load for
154 CGNS/HDF5 is available in CHLone for example. CHLone (<http://chlone.sourceforge.net>)

155 is a python module implementing the SIDS-to-Python mapping which allows a simple
156 load/save of CGNS/HDF5 files.

157 3. direct access to a face connectivity

158 For some specific unstructured software, it is interesting to have a simple access to a specific
159 face in the ElementConnectivity array.

160 This is only allowed with the “Face Position” solution.

161 4. Access to face position and face number of nodes

162 Using the incremental index method, it is easy to calculate the “face number of nodes” for
163 element ‘i’ as ElementStartIndex[i+1]-ElementStartIndex[i] is an O(1) calculation. So the
164 selected method gives both data (face position and face number of nodes) in O(1) where the
165 alternate method would give face number of nodes in O(1), but face position in O(N).

166 **2.2.5 ElementStartIndex content**

167 Using an incremental index --the face position in the ElementConnectivity array-- could lead
168 to large numbers in the ElementStartIndex array. However, these numbers are limited to the
169 size of the ElementConnectivity array as they give access to addresses of that array. We will
170 probably reach the ElementConnectivity array limits first.

171 **2.2.6 ElementStartIndex last value**

172 The ElementStartIndex lists the first vertex position of each face in the
173 ElementConnectivity. Then we added the ElementConnectivity size as last value. The
174 rationale behind this is to improve data access when looping over the faces.

175 → Positions without ElementConnectivity size

176 In this configuration an iteration loop over NGONs could be written as follows:

```
177 for (i=0; i < NGON; i++) {  
178     end = (I == NGON-1) ? ElementConnectivitySize : ElementStartIndex[i+1];  
179     for (j=ElementStartIndex[i]; j < end; j++) {  
180         entry = ElementConnectivity[j];  
181         // do something here  
182     }  
183 }
```

184 For each NGON we have to check if we are reaching the last NGON to handle the access to
185 the last vertex of the last face.

186 → Positions with ElementConnectivity size (current proposal)

187 In this configuration an iteration loop over NGONs could be written as follow:

```
188 for (i=0; i < NGON; i++) {  
189     for (j=ElementStartIndex[i]; j < ElementStartIndex[i+1]; j++) {  
190         entry = ElementConnectivity[j];  
191         // do something here  
192     }  
193 }
```

194 The last value of the ElementStartIndex gives a boundary to access the last vertex of the last
195 NGON without the conditional assignment.




Illustration 6: Simple configuration demonstrating the NGON proposal.

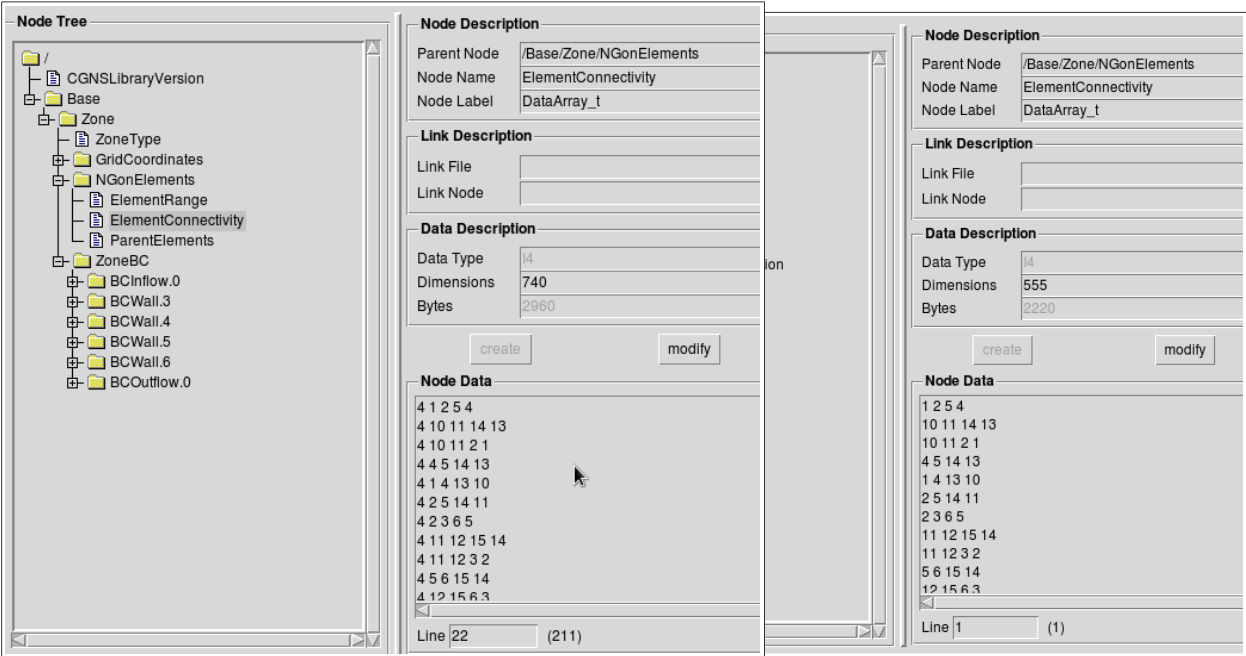


Illustration 7: Previous vs. new ElementConnectivity array

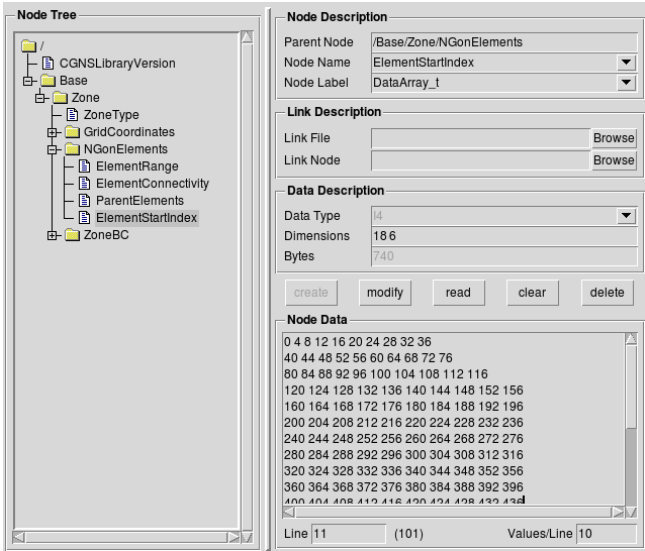


Illustration 8: ElementStartIndex array

197 4 Implementation note

198 This section aims to reflect discussions during previous CGNS meetings related to the
199 implementation impact of this modification proposal. As detailed above, this proposal strongly
200 modifies the representation of the ElementConnectivity node. As a consequence we need to address
201 the backward compatibility issue.

202 The following solutions were discussed in order to insure backward compatibility:

- 203 • CGNSLibraryVersion
204 In software, the CGNSLibraryVersion could be used to determine whether the CGNS file
205 uses the current or proposed NGON representation. This number attached to the standard
206 version will allow parsing tools to use the correct NGON representation.
- 207 • Conversion utility
208 A dedicated tool could be written to convert current CGNS files to the new NGON
209 representation
-

210 These two solutions should allow a smooth transition for all users of the CGNS standard.

211 5 Conclusions

212 This proposal for the NGON representation addresses the HPC issue due to the interlaced data
213 representation of the unstructured connectivity description. The solution optimizes the data
214 representation for parallel IO and has a low impact on the SIDS to deinterlace data.

215 As shown in the first part, MIXED and NFACES elements are equally concerned by the parallel IO
216 access. As such their SIDS representation should be updated accordingly to the solution chosen for
217 NGON. This issue should be addressed in a later CPEX.

218 6 Appendix: Document modification list

219 1. Following Marc Poinot remarks:

- 220 • Rename array from FaceConnectivityPosition to ElementStartIndex (text and figures)
- 221 • Reformulate a sentence describing the ElementStartIndex at line 101
- 222 • Remove remark concerning array size limitation in section 2.1.2.5

223 2. Following Robert Bush suggestion:

- 224 • Added section “Implementation note” to reflect CGNS committee discussions.

225 3. Remove multiple typo.

226 4. Following Gregory Sjaardema feed back:

- 227 • Modify the ElementStartIndex definition to improve the ability to loop on the data array.
228 The new size is N+1 and last index represent the ElementConnectivity total size.
229 This modification is propagated through sections 2 and 3.
- 230 • Correct multiple grammar, typographical or formatting issues.